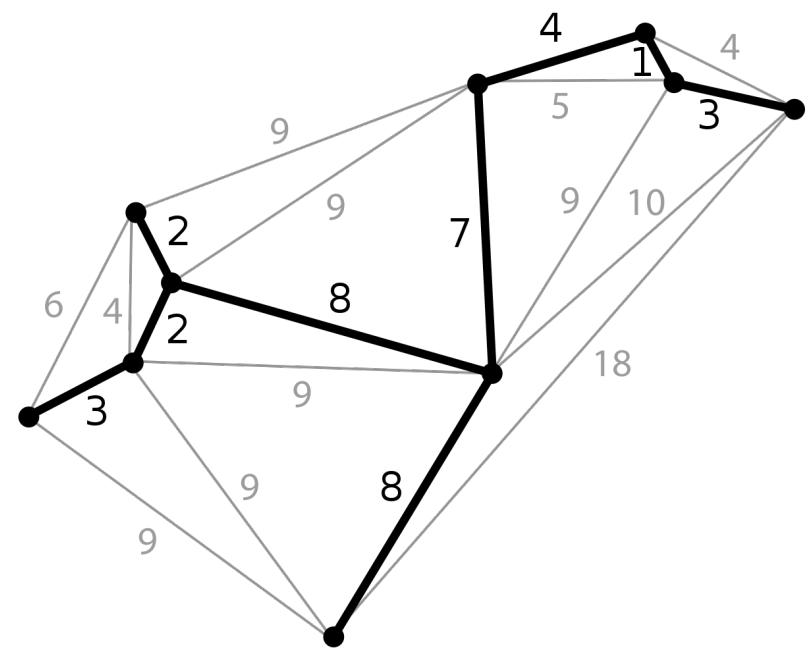


# Data Structures and Algorithms (DSA) for AI

Katja Tuma



# Single-source shortest path with Dijkstra

## Problem formulation

**Input:** directed  $G = (V, E)$ ,  $s \in V$ , non-negative length  $l_e$  for each edge  $e \in E$

**Output:**  $dist(s, v)$ , for all  $v \in V$

# Single-source shortest path with Dijkstra

## Problem formulation

**Input:** directed  $G = (V, E)$ ,  $s \in V$ , non-negative length  $l_e$   
for each edge  $e \in E$

**Output:**  $dist(s, v)$ , for all  $v \in V$

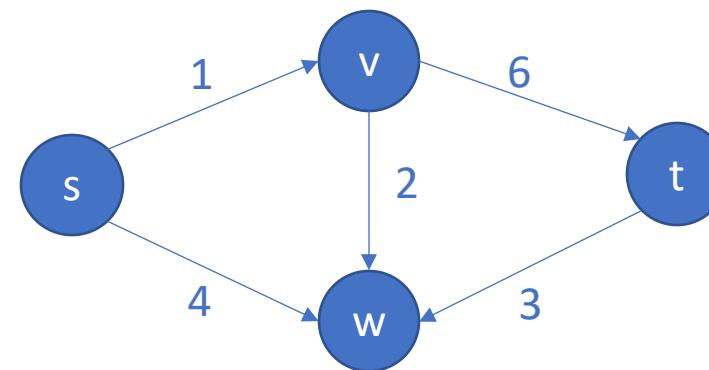
Dijkstra assumptions: 1) directed, 2) non-negative edge length

Additional assumption: there exist paths from  $s$  to every vertex  $v$  in  $V$ .

Main idea: Explore the cheapest next edge (according to the Dijkstra score)

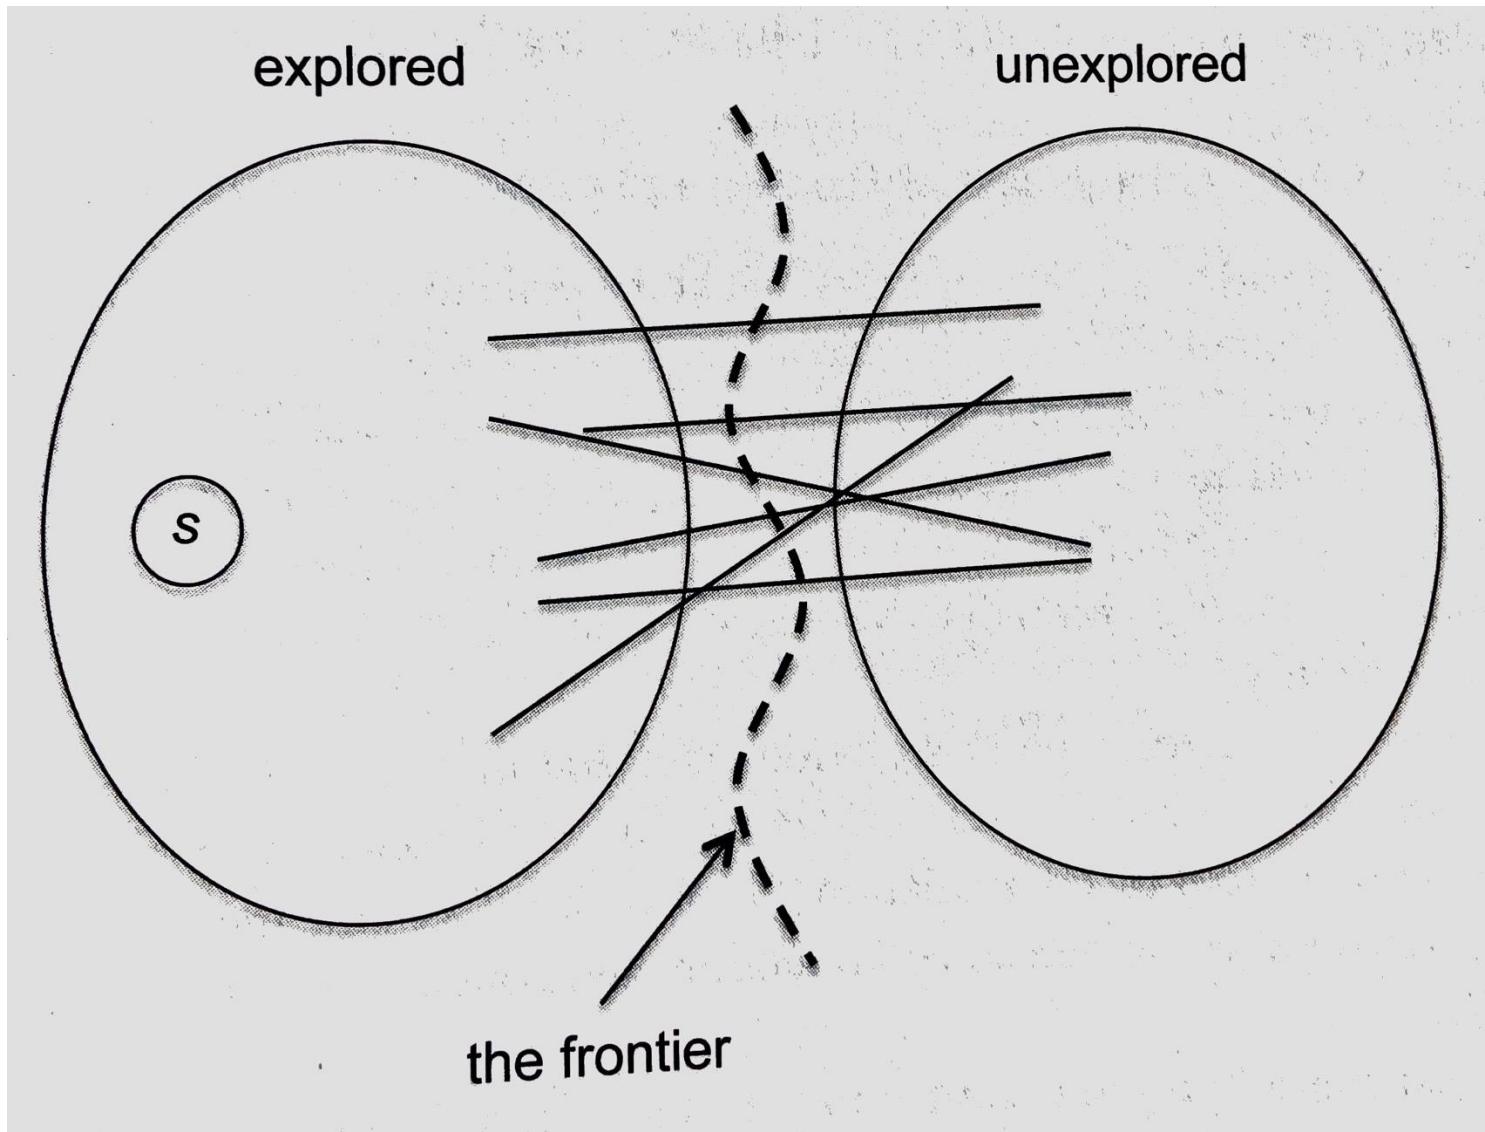
# Single-source shortest path with Dijkstra

- $X$  explored
- $d_{score} = \text{len}(v) + l_{vw}$
- $V - X$  unexplored
- starting vertex =  $s$



How to  
choose an  
edge?

$$d_{score} = \text{len}(v) + l_{vw}$$



# Dijkstra - pseudocode

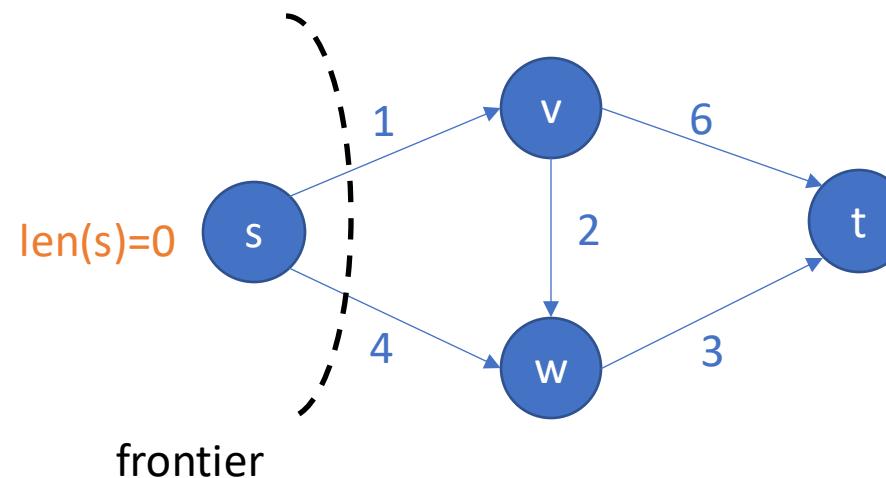
**Input:** directed  $G = (V, E)$  in adjacency list representation,  $s \in V$ , non-negative length  $l_e$  for each edge  $e \in E$

**Postcondition:** for every  $v \in V$  the value  $\text{len}(v)$  is the true shortest path distance  $\text{dist}(s, v)$ .

```
X = {s}                                //Initialization
len(s) = 0, len(v) = +∞ for all v≠s    //Initialization
//main loop
while there is an edge (v, w), v ∈ X and w ∉ X do:
    (a, b) = such an edge minimizing len(v) + lvw
    add b to X
    len(b) = len(a) + lab
```

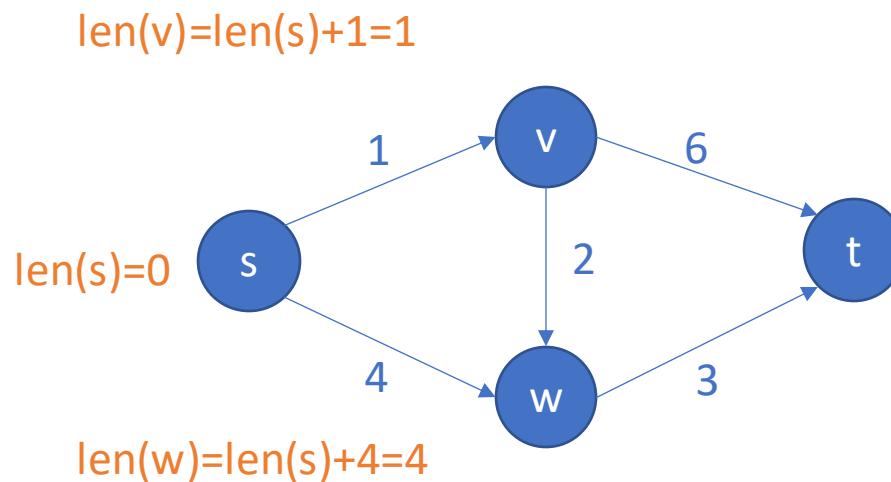
# Single-source shortest path with Dijkstra

- $d_{score} = \text{len}(v) + l_{vw}$



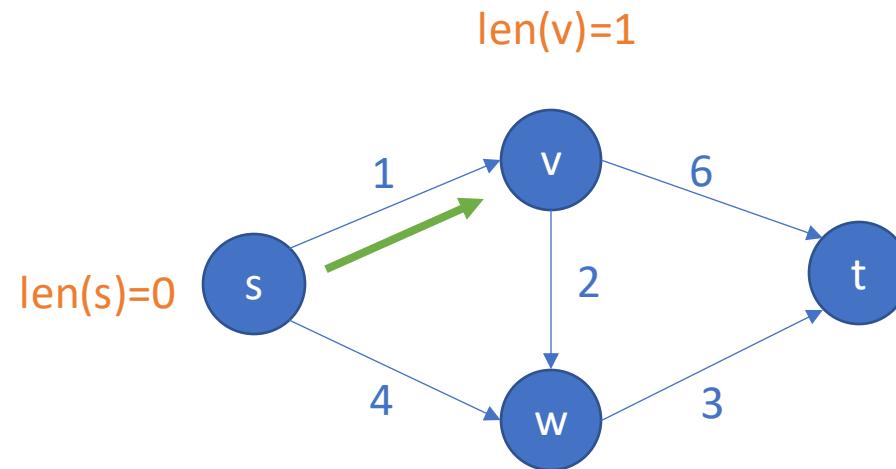
# Single-source shortest path with Dijkstra

- $d_{score} = \text{len}(v) + l_{vw}$



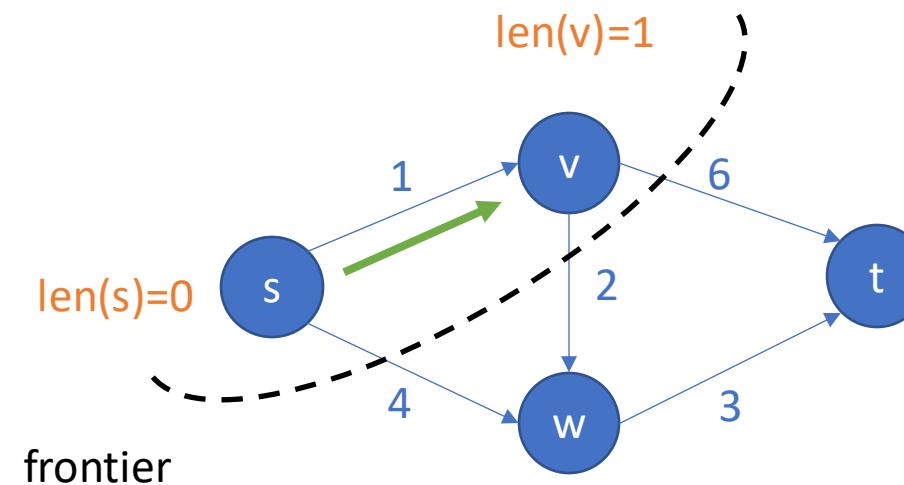
# Single-source shortest path with Dijkstra

- $d_{score} = \text{len}(v) + l_{vw}$



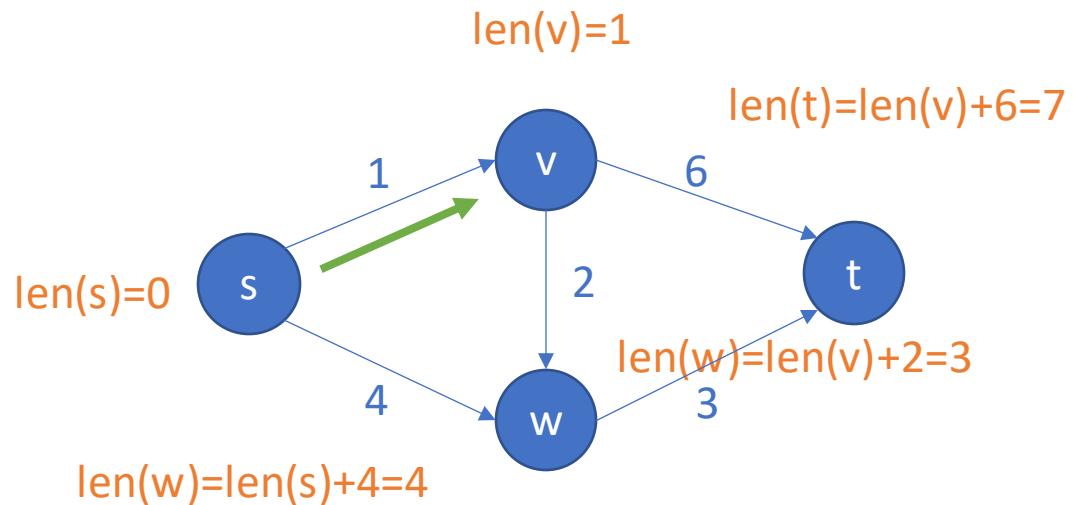
# Single-source shortest path with Dijkstra

- $d_{score} = \text{len}(v) + l_{vw}$



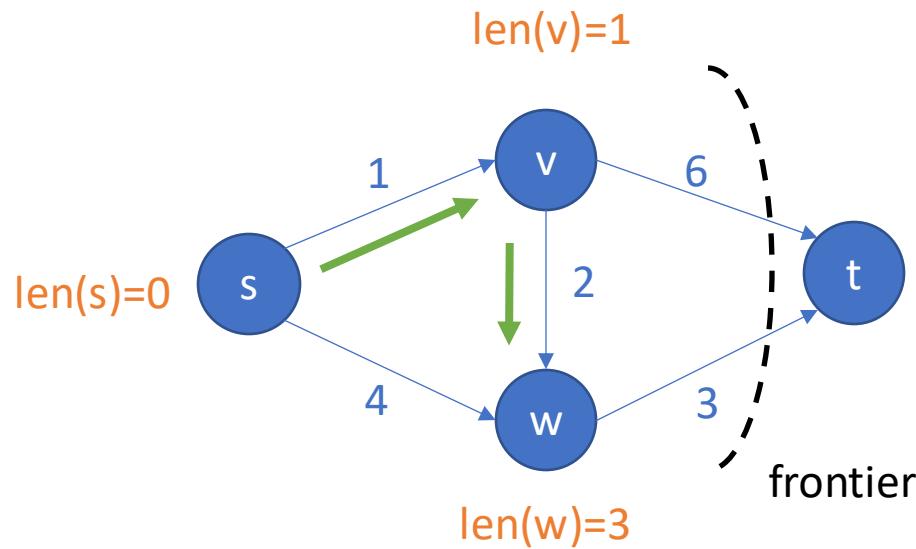
# Single-source shortest path with Dijkstra

- $d_{score} = \text{len}(v) + l_{vw}$



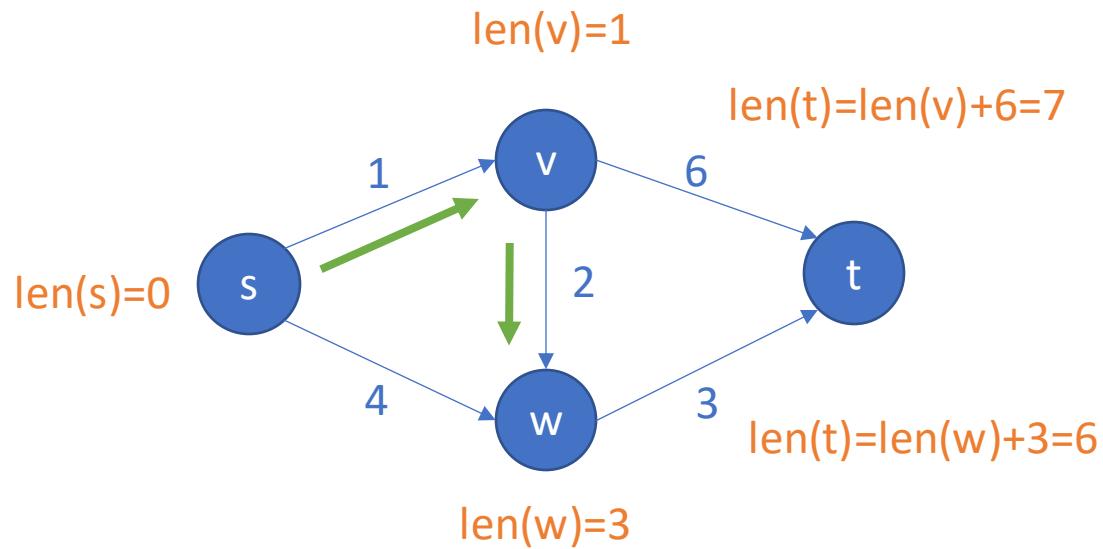
# Single-source shortest path with Dijkstra

- $d_{score} = \text{len}(v) + l_{vw}$



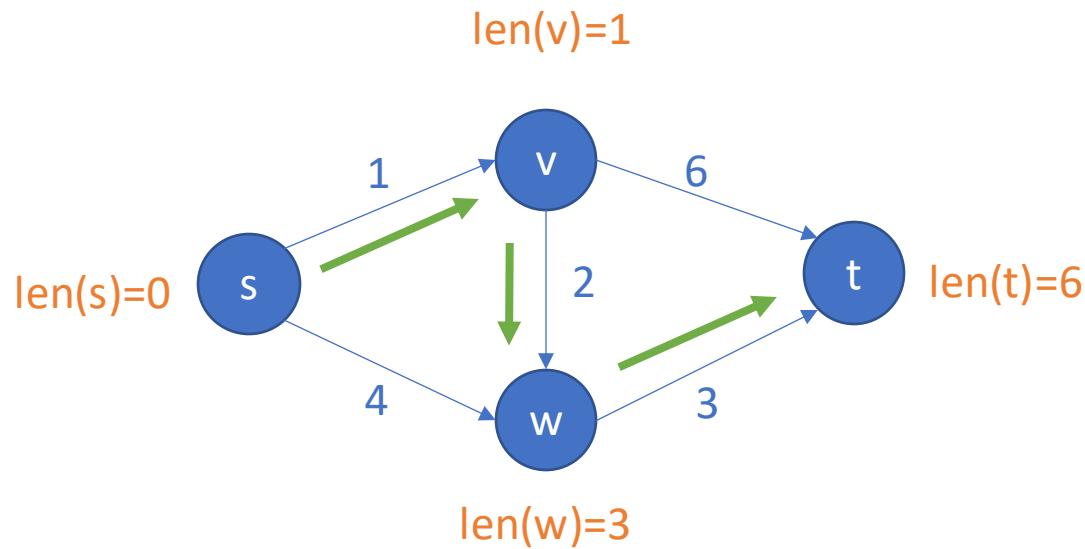
# Single-source shortest path with Dijkstra

- $d_{score} = \text{len}(v) + l_{vw}$



# Single-source shortest path with Dijkstra

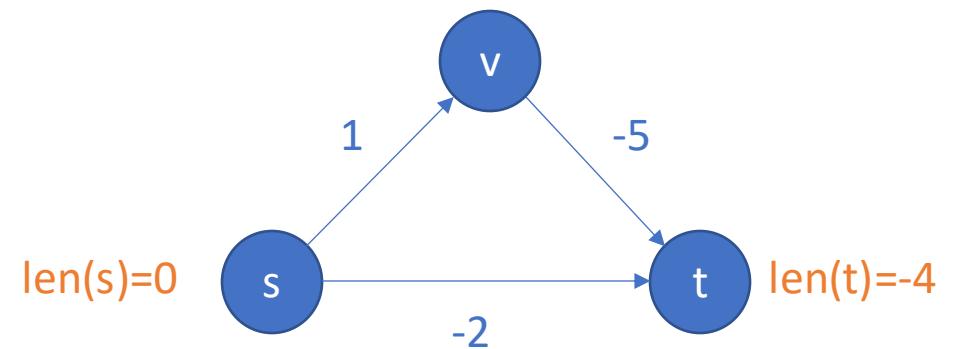
- $d_{score} = \text{len}(v) + l_{vw}$



# Why not negative weights?

A bad example for Dijkstra:

- Two possible paths  $s \rightarrow t$ :
  - $s-v-t$  with cost  $0+1+(-5)=-4$
  - $s-t$  with cost  $0+(-2)=-2$
- Since  $-4 < -2$ , the true shortest path is  $-4$

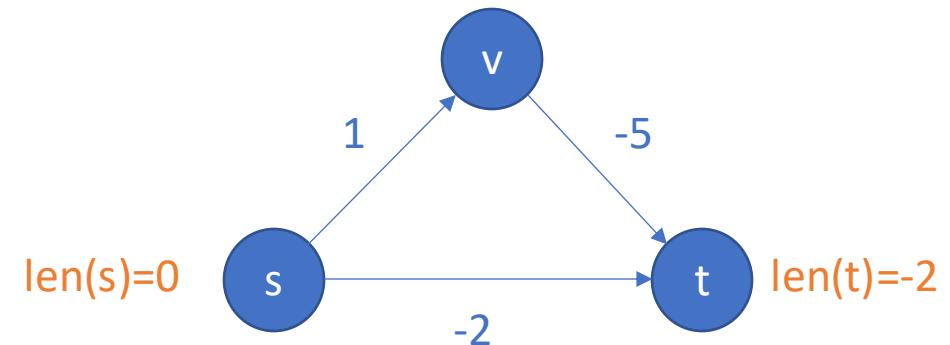


- What happens if we receive this graphs as input?

# Why not negative weights?

What happens if we receive this graphs as input?

- $X=\{s\}$  and  $\text{len}(s)=0$
  - First iteration:
    - Dijkstra Scores:
      - $(s,v) = 0+1 = 1$
      - $(s,t) = 0+(-2) = -2$
    - Concludes that  $t$  is the cheapest edge.
    - Add  $t$  in  $X$
  - Returns the shortest path is  $s-t$ .
  - Incorrect



# Why not negative weights?

Can we reduce a problem to something we already know how to solve?

Why not simply add a large enough positive number to all edges?

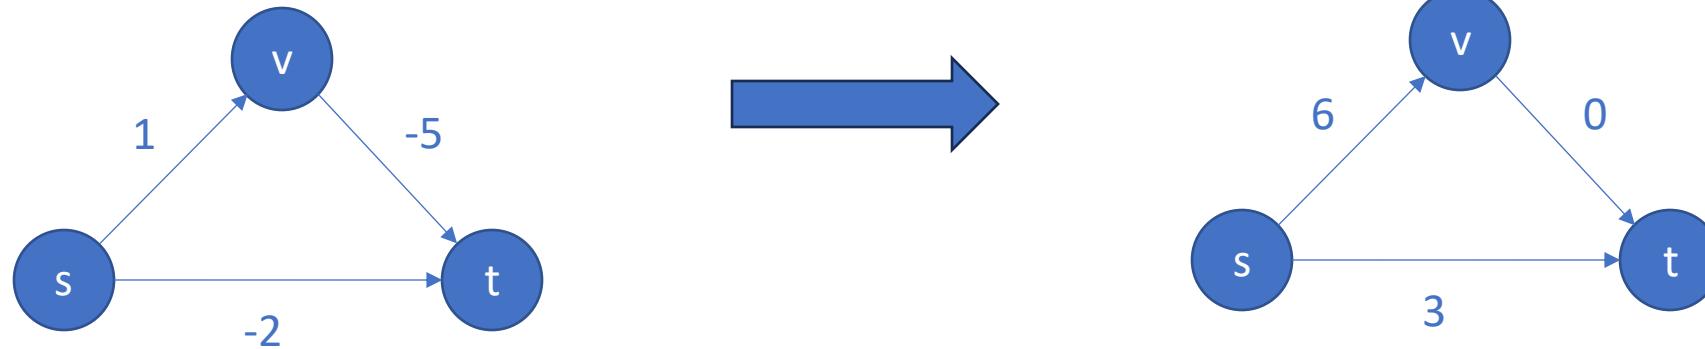
Eg +5

# Why not negative weights?

Can we reduce a problem to something we already know how to solve?

Why not simply add a large enough positive number to all edges?

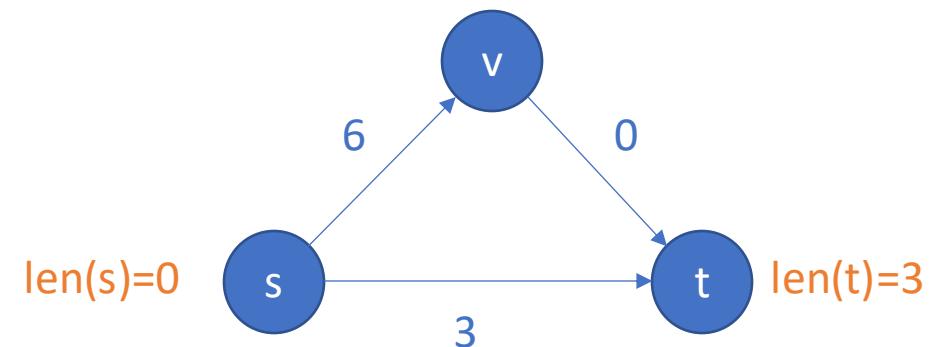
Eg +5



# Why not negative weights?

What happens if we receive this graphs as input?

- $X=\{s\}$  and  $\text{len}(s)=0$
  - First iteration:
    - Dijkstra Scores:
      - $(s,v) = 0+6 = 6$
      - $(s,t) = 0+3 = 3$
    - Concludes that  $t$  is the cheapest edge.
    - Add  $t$  in  $X$
  - Returns the shortest path is  $s-t$ .
  - Incorrect



# Correctness of Dijkstra

- Proof by induction in the book  
 $P(1)$  ...shortest path to  $s=0$  (base case)

$P(k)$  set the inductive hypothesis and assume  $P(k)$  is true

$P(k+1)$  use the IH and the assumption about non-negative edge lengths to reason that the Dijkstra computes true shortest paths for any vertex in the graph

- Online: <http://www.algorithmsilluminated.org>

# Dijkstra - pseudocode

**Input:** directed  $G = (V, E)$  in adjacency list representation,  $s \in V$ , non-negative length  $l_e$  for each edge  $e \in E$

**Postcondition:** for every  $v \in V$  the value  $\text{len}(v)$  is the true shortest path distance  $\text{dist}(s, v)$ .

$X = \{s\}$

$\text{len}(s) = 0$ ,  $\text{len}(v) = +\infty$  for all  $v \neq s$

//main loop

while there is an edge  $(v, w)$ ,  $v \in X$  and  $w \notin X$  do:

$(a, b) =$  such an edge minimizing  $\text{len}(v) + l_{vw}$

add  $b$  to  $X$

$\text{len}(b) = \text{len}(a) + l_{ab}$

Can we do better?

# How to improve?

- Minimum heap can help

# Invariant

---

The key of a vertex  $w \in (V - X)$  is the minimum Dijkstra score of an edge with tail  $v \in X$  and head  $w$  OR  $+\infty$  if no such edge exists.

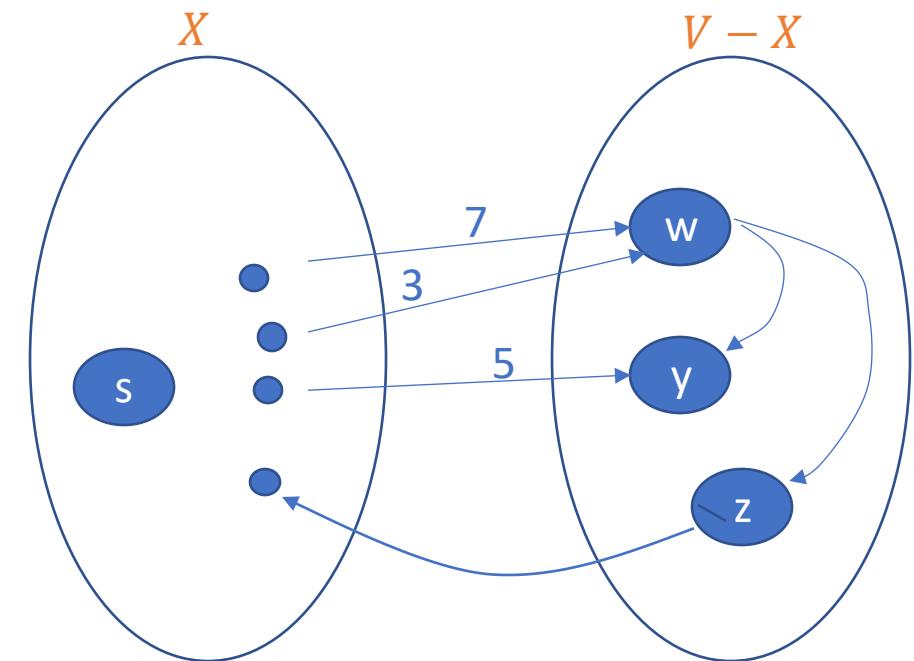
# How choosing the next edge works?

- $key(w) = \min_{\substack{(v,w) \in E, \\ v \in X}} (len(v) + l_{vw})$

- Two-round knockout tournament
  - 1. ‘local’: 3 vs 7: 3 wins
  - 2. ‘global’: 3 vs 5: 3 wins

→ $key(w)=3$

→Extract min from heap & add the next vertex (w) to X (frontier changes)



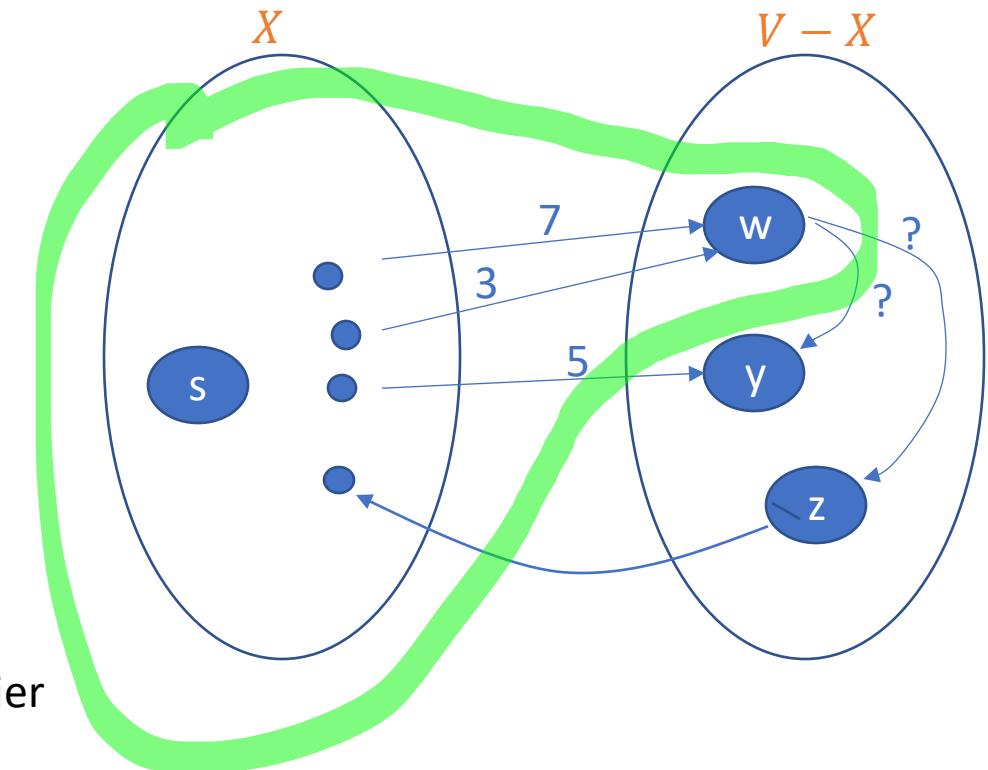
# Dijkstra with heap – pseudocode

```
X = empty set; H = empty set
key(s) = 0
for every  $v \neq s$  do
    key( $v$ ) =  $+\infty$ 
//use heapify (insert all other vertices into H)
while H is not empty do
     $w$  = extractmin( $H$ )
    add  $w$  to  $X$ 
    len( $w$ ) = key( $w$ )
    //UPDATE HEAP...pay the price
```

# Maintaining the invariant

INVARIANT: The key of a vertex  $w \in (V - X)$  is the minimum Dijkstra score of an edge with tail  $v \in X$  and head  $w$  OR  $+\infty$  if no such edge exists.

- We have to maintain the invariant, ie fix the heap (recalculate '?' key( $y$ )=5 and key( $z$ )= $+\infty$  need to be updated)
- There are new edges on the frontier
- They need to get their key recalculated in the heap



# Dijkstra with heap – pseudocode

```
X = empty set; H = empty set
key(s) = 0
for every  $v \neq s$  do
    key( $v$ ) =  $+\infty$ 
//use heapify (insert all other vertices into H)
while H is not empty do
    w = extractmin(H)
    add w to X
    len(w) = key(w)
    //UPDATE HEAP...pay the price
        ↗ //UPDATE HEAP
        foreach edge (w, y) with  $y \in V - X$  do:
            DELETE y from H
            key(y) = min(key(y), len(w) + lvw)
            INSERT y into H
```

# Running time analysis

You check: the running time is dominated by the heap operations! (HW)

What work is done for heap ops?

- (n-1) Extract mins (which triggers the heap update = delete+insert)

How many delete/insert?

- a vertex can have as many as n-1 outgoing edges (scary! That would mean  $n^2$  heap operations). True for DENSE graphs

→ i.e., many “local tournaments”

- in general, much better. Remember we **only update the key if the tail vertex has been sucked into X.**
- **each edge only triggers at most one** Delete / Insert combo (if v added to X first)

So: # of heap operations is  $O((n-1)+m)=O(n+m)$ . Since we assumed that there exist all paths (from s to any v), ie the graph is weakly connected, we know that m dominates n. So, we can simplify  $O(m)$ .

So: running time =  $O((m+n)\log n)$  OR, simplified under the assumption  $O(m * \log(n))$ .